

Persons Required: Context & Framework

AI 201 — CONTEXT & FRAMEWORK

Project 3: Persons Required

The research process, the ESF practices, the user testing playbook, and the case study format.

The Shift

Projects 1 and 2 asked you to direct AI as an Art Director building screens and systems. You wrote visual specs. You evaluated AI output against color palettes and hover states. You got good at it. Fast.

Project 3 is different. The spec you write before AI engagement is no longer a visual document. It is a design argument: a thesis about a real person, a real problem, and what “helped” looks like. The visual decisions still matter — but they serve the argument, not the other way around.

This is the shift from **designer** to **design thinker**. You are no longer proving you can build. You are proving you can see a problem, understand it deeply enough to have a position on it, and then build something that addresses it.

The AI tools do not change. Your relationship to them does. When you direct AI in this project, you are directing it on behalf of someone who is counting on you to get it right.

Milestone Schedule

Session	Date	Milestone
Session 13	Mon 5/4	The Room Debate. Who needs help and why do you care? Thesis pitch to the room.
Session 14	Wed 5/6	Research & Field Work. Interview your person. Document the problem. Platform decision defended. Design Argument due.

Session 15	Mon 5/11	Build begins. Design Intent locked. AI translates your spec into a working prototype.
Session 16	Wed 5/13	★ FIRST CONTACT. Put the prototype in front of your person. Document what happens. Not a crit — a field test.
Session 17	Mon 5/18	Iteration. Rebuild based on what you learned. Second test if possible. Collect evidence.
Session 18	Wed 5/20	The Case Study. Build the story: the person, the problem, the prototype, the evidence. QA and polish.
Session 19	Mon 5/25	Final Polish. Live URL verified. Git history clean. Case study presentation rehearsed.
Session 20	Wed 5/27	★ PREMIERE + STUDIO CRIT. Present the case study. Not a demo — a defense. The person, the problem, the proof.

The Design Argument

In Projects 1 and 2, your Design Intent was a visual spec: palette, typography, hover behavior, mood. It told AI what the thing should look like.

In Project 3, your pre-AI document is a Design Argument. It tells AI — and yourself — why this thing should exist.

What a Design Argument Contains

The Person. Not a persona. A real human being with a name, a context, and a life you have directly observed. Who are they? What is their daily reality? What do they deal with that they shouldn't have to?

The Problem. What is broken, painful, confusing, or missing? Be specific. "My grandmother has trouble with technology" is not a problem statement. "My grandmother misses her medication because her phone's reminder system requires six taps to dismiss and she panics and closes the app" is a problem statement.

The Definition of "Helped." What does success look like for this person? Not for your portfolio. For them. How will you know it worked? This is measurable or observable — not aspirational.

Your Qualification. Why are you the right person to attempt this? What do you know, understand, or have access to that makes your perspective valuable? This is where your discipline, your experience, and your relationship to this person become your unfair advantage.

The Platform Decision. Where does this thing need to live? A web app? A phone? A physical space? A browser extension? The person and the problem determine the platform — not your comfort zone.

Non-Negotiables. What must this project do or be, no matter what? What will you not compromise on, even if AI suggests a faster or easier alternative? These should be driven by what your person needs, not what looks good in a demo.

The Design Argument is written before AI engagement.

This is non-negotiable. If AI helped you identify the problem, AI is holding the intellectual authority, not you. You must know what you believe before you ask AI to help you build it. The Design Argument is the standard against which you evaluate everything AI produces for this project.

The Research Process

This is not a UX research course. You are not expected to run a formal study. But you are expected to do the foundational work of actually understanding the person you are designing for. This is the Industrial Design tradition: go where the problem lives, observe, listen, then come back and build.

Step 1: Choose Your Person

This must be someone you can access directly. You will need to talk to them at least twice (once for research, once for testing) and ideally more. Family members, friends, coworkers, community members, mentors, neighbors. Someone you care about enough to do this well.

Ethical Note

You are not conducting formal human subjects research, so IRB approval is not required. However, you must: (1) Get verbal permission before interviewing or observing. (2) Get permission before photographing or recording. (3) Let your person know their problem and feedback will appear in a

class presentation. (4) Use first names only in documentation unless they explicitly consent to more. Respect is the baseline.

Step 2: Listen Before You Solve

The most common mistake in design-for-real-people projects is arriving with a solution and looking for a problem to justify it. Do not do this.

Your first conversation with your person is not about your project. It is about their life. Ask open questions. Let them talk. Take notes. Pay attention to the things they mention casually that sound painful — those are often the real problems, not the ones they'd formally complain about.

Questions That Work

Ask This	Not This
Walk me through your morning routine.	What apps do you wish existed?
What's the most frustrating part of [activity]?	Would you use an app that does X?
Show me how you currently do [task].	What features would you want?
What do you wish someone understood about your [situation]?	Would this solve your problem? [shows wireframe]
When was the last time [problem] really got to you?	Rate your frustration on a scale of 1–10.

The left column gets you stories. The right column gets you yes/no answers that confirm your assumptions instead of challenging them. Stories are where the design insights live.

Step 3: Document the Context

After your conversation, document what you observed. This is your research documentation deliverable. Include:

Direct quotes — the person's actual words. These are more valuable than your interpretation.

Environment details — where does this problem happen? What does the space look like? What tools are already present? Photos with permission.

Workarounds — how is the person currently coping? Sticky notes, phone alarms, asking someone else, avoidance? Workarounds reveal what the system should replace.

Constraints — what can't change? Budget, physical ability, tech literacy, available devices, internet access, time, environment. These are your design constraints and they are gifts, not obstacles.

User Testing: First Contact

Session 16 is called First Contact for a reason. This is the moment your project leaves the screen and enters the real world. You are putting an unfinished prototype in front of the person you designed it for and watching what happens.

This is not a demo. This is not a crit. This is a field test.

How to Run a First Contact Test

1. Set expectations. Tell your person: "This is rough. Things will be broken. I'm not showing you a finished product — I'm showing you a direction. I need to see where it's right and where it's wrong."

2. Shut up. Give them the device or the URL. Do not explain how it works. Do not point at buttons. Do not say "you're supposed to tap that." Watch. If they get stuck, note where. If they do something unexpected, note what. Their confusion is data.

3. Ask what happened. After they've used it: "What did you expect to happen?" "What was confusing?" "What would you change?" "Would you use this tomorrow?" That last question is the one that matters.

4. Document everything. Photos of them using it. Screen recordings if they consent. Their exact words. What they reached for that wasn't there. What they ignored that you thought was important. This is your User Testing Evidence deliverable.

The hardest part of user testing is hearing that your design is wrong.

When your person struggles with something you built, the instinct is to explain why they're using it incorrectly. Fight that instinct. If they're confused, the design is confused. That's not failure — that's the entire reason you're testing before you ship. The best Records of Resistance in this project will come from this moment.

The Case Study Presentation

Session 20 is not a demo day. It is a case study defense. You are presenting the full arc: the person, the problem, the research, the build, the test, the iteration, and the evidence.

Case Study Structure

Section	What You Show

The Person	Who are they? Why do they matter to you? One photo (with permission) or description. Set the human context before anything technical.
The Problem	What is broken? Show your research: quotes, environment photos, the workarounds they were using before. Make the room feel the problem.
The Thesis	Your Design Argument in one sentence. What does “helped” look like? Why this platform?
The Prototype	Live demo of the product. Walk through the core flow. Show the design decisions that came directly from your research.
First Contact	What happened when your person used it? Show the evidence: photos, recordings, quotes. What worked? What failed? Be honest.
The Iteration	What did you change based on testing? Show before/after. Explain why. This is where your editorial judgment is most visible.
The Evidence	Did it help? Show the proof. Quotes from your person. Observed behavior changes. Screenshots of real use. If it didn't fully work, say so and explain what you'd do next.

A strong case study makes the room care about the person before they see the product. If you show the product first, it's a demo. If you show the person first, it's a story.

Portfolio Tip

This case study format is the industry standard for UX/product portfolios. A project that ships with research, testing evidence, and iteration documentation is worth ten projects that only have screenshots. If you're building your book, this is the one that goes on the front page.

ESF Practices in Project 3

The Epistemic Stewardship practices are the same. The stakes are higher because there's a real person on the other end of your decisions.

The Five Questions

Question	What It Means in Project 3
Can I defend this?	Can you explain every design decision by pointing back to your research and your person's needs?
Is this mine?	Did you direct AI based on your Design Argument, or did you accept AI's suggestions because they looked good? When your person is counting on you, "it looked cool" is not a defense.
Did I verify?	Does the product actually work the way your person needs it to? Did you test with them, not just in your browser?
Would I teach this?	Could you explain the system architecture, the research process, and the design rationale to another designer? Do you understand it well enough to hand it off?
Is my disclosure honest?	Does your AI Direction Log accurately reflect what happened? When you're presenting a case study with real evidence, fabrication is not just academic dishonesty — it's professional fraud.

Records of Resistance

In Projects 1 and 2, your Records of Resistance were about visual decisions: AI chose the wrong color, the wrong layout, the wrong animation. Those still matter in Project 3.

But the most important resistance moments in this project will be about the person. AI does not know your person. It will generalize. It will suggest solutions based on patterns, not on the specific human being sitting across from you. When AI gives you a generic "accessibility dashboard" and your person needs something completely different, that rejection is a Record of Resistance — and it's the most valuable one you'll write all quarter.

AI Direction Log

Five or more entries for Project 3. The additional entries reflect the additional complexity: you're now directing AI across research synthesis, architecture decisions, platform-specific implementation, and iteration based on real user feedback. Your log should capture decisions across the full arc, not just the code phase.

The Platform Decision

React is our shared classroom language. It is not required for Project 3. The person you are helping determines where your product needs to live.

Platform Decision Framework

If Your Person Needs...	Consider...
Something on their phone	React (mobile-responsive PWA), React Native, or a native app framework. A PWA is the fastest path if the person just needs a mobile-friendly URL.
Something in their browser workflow	Chrome extension, bookmarklet, or browser-based tool. If the problem lives inside Chrome, the solution should too.
Something spatial or immersive	Unity (WebGL or standalone build) or Unreal. If the person needs to move through a space or interact with 3D objects, a web page won't do it.
Something in a physical space	Kiosk mode web app, projected display, Raspberry Pi installation, or printed material generated by code. The platform might not be a screen at all.
Something inside a tool they already use	Figma plugin, Slack bot, Discord bot, Google Sheets add-on, VS Code extension. Meeting people where they already work is often better design than building a new destination.
Something with AI inside it	React or any platform + API integration (Claude API, Gemini API). If AI is a feature of the product, not just a build tool, discuss the architecture with the instructor during Session 13.

If you choose a non-React platform, **discuss your plan with the instructor during Session 13.** Your Design Argument must include the platform rationale. The ESF practices are identical across all pipelines.

Common Pitfalls

Solving your own problem instead of someone else's. Building a tool for yourself is tempting. But you already know your own workarounds, tolerances, and preferences. Designing for someone else forces you to listen, test assumptions, and build empathy into the process. That's the whole point.

Choosing a person you can't access. If you can't sit across from this person twice in four weeks, choose someone else. Access is a hard constraint. Plan for it.

Scoping too big. You are not building a startup. You are building one tool that solves one problem for one person. If your Design Argument requires three paragraphs to explain the problem, the scope is too wide. Narrow it until you can say it in one sentence.

Skipping the test. The user testing is not optional. A product that has never been in front of its intended user is an assumption, not a solution. If your person can't test in person, a video call with screen share counts. No test, no evidence, no credit.

Letting AI define the problem. AI is excellent at generating plausible problem statements. Do not use it for this. Your understanding of the problem must come from your person, not from a prompt. If Claude wrote your thesis, the Five Questions will catch it.

The Art Director in Project 3

You have been the Art Director all quarter. In Projects 1 and 2, the Art Director's authority came from a visual spec: you knew what the thing should look like, so you could evaluate AI's output against that standard.

In Project 3, your authority comes from something deeper: you know the person. You've sat with them. You've heard their story. You understand the problem in a way that AI never will, because AI has never been scared in a hospital bed, or frustrated by a government form, or lost trying to navigate a system that wasn't designed for someone like them.

That understanding is your Design Intent. It is the standard against which you evaluate every line of code AI writes, every layout it proposes, every feature it suggests. When AI drifts from what your person needs, you pull it back. That's not resistance for the sake of resistance. That's advocacy.

Academic Policies

AI Use

AI tools (Claude, Gemini, and any others you choose) are permitted and expected for production

tasks: code generation, layout, debugging, data formatting, responsive design, and similar build work. AI is your engineering team.

AI is not permitted for research tasks: the interview, the observation, the problem identification, and the Design Argument thesis must be your own work, based on your direct experience with your person. AI may not write your Design Argument, your research documentation, or your user testing notes. AI may clean up grammar or formatting in these documents after you have written them.

Your AI Direction Log and Records of Resistance must be honest and specific. The Five Questions reflection is your final self-audit. Fabrication in any documentation is an academic integrity violation.

Attendance

Sessions 13 (Room Debate) and 20 (Premiere) require in-person attendance. If you cannot attend Session 20, contact the instructor before Session 19 to arrange an alternative presentation time. Unexcused absence from Session 20 results in a zero for the Case Study + Post-Mortem category (10 points).

Late Work

Project 3 is due at the start of Session 20 (Friday, May 29, 2026). Late submissions lose 10 points per calendar day. Submissions more than 3 days late are not accepted without prior arrangement.

■ DIRECTIVE MEMO (Human-Authored)

This assignment exists because design is not decoration. It is an act of seeing someone clearly, understanding their problem deeply, and building something that makes their life measurably better. The technical skills from Projects 1 and 2 are the means. The person is the end. If the portfolio piece is stunning and the person was not helped, the project failed. If the portfolio piece is rough and the person was genuinely helped, the project succeeded. — T. Lindsey